

Kernel Principal Component Analysis

Alex Beeny abeeny@siue.edu

Draft: April 2, 2024

Contents

1	Introduction	2
2	Principal Component Analysis	3
2.1	Finding uncorrelated features	3
2.2	Minimizing projection residuals	4
2.3	Singular value decomposition	5
2.4	Principal component analysis algorithm	5
2.5	Linear regression and PCA	7
3	Reproducing Kernel Hilbert Space	8
3.1	Kernel methods	12
3.2	Constructing kernels	13
4	Kernel PCA	16
4.1	Covariance matrix and kernel matrix	16
4.2	Centering in the feature space	16
5	Conclusion	17
A	Linear Algebra	17
A.1	Matrix operations and notation	19
A.2	Broadcasting	20
B	Riesz Representation Theorem	20
C	Mercer's Theorem	21
D	Code	21

Abstract

Principal component analysis (PCA) and kernel methods are tools often used in data science. The underlying theory of these tools depend on the properties of a special type of Hilbert space called a reproducing kernel Hilbert space (RKHS). This paper explores the essence of RKHSs




Figure 1: Plotting y against x (left) and the derived feature $z = \phi(x)$ (middle). The 3D plot (right) graphs the output y against x and x^2

using data science examples, in particular, PCA and kernel PCA. When kernel methods are applied to PCA, we can analyze nonlinear data in a high-dimensional feature space with some nice properties.

1 Introduction

In linear regression, the equation of a line $y = a_0 + a_1x$ is used to model observations based on training data. Here, the input variable x is used to predict the response variable y . The parameters a_0 and a_1 are chosen such that the residual error¹ is minimized. It seems natural to model data using polynomial equations in a similar way, that is, determine a_0, a_1, \dots, a_n such that

$$y = a_0 + a_1x + a_2x^2 + \dots + a_nx^n \quad (1)$$

minimizes the residual error.

In multiple linear regression, the equation of a hyperplane

$$y = a_0 + a_1x_1 + a_2x_2 + \dots + a_nx_n \quad (2)$$

is used to predict y using inputs x_1, x_2, \dots, x_n . It follows that these observations are points in $n + 1$ dimensions.

Example 1.1. Suppose we have a set of points $\{(x^{(i)}, y^{(i)})\}_{i=1}^k$ in \mathbb{R}^2 . Using polynomial regression, we fit the model $y \sim 1 + x + x^2$. We can derive a new variable from x to get

$$z = \phi(x) = a_0 + a_1x + a_2x^2.$$

By transforming our original data, the quadratic relationship between x and y can be viewed as a linear relationship between z and y . This shows that polynomial regression is just a special case of multiple regression where each power of x is treated as a separate dimension. See Figure 1

Here, we make the distinction between different kinds of input variables. We say that x is an **attribute** and that z is a **feature**.

¹The residual for a given observation $(x^{(i)}, y^{(i)})$ is $|y^{(i)} - (a_0 + a_1x^{(i)})|$.

2 Principal Component Analysis

When analyzing data, it can be convenient to transform the given input variables to produce new features. For a well-chosen transform, these features may be approximated using fewer dimensions than the original input space [6]. This is an example of a data preprocessing technique known as *dimension reduction* and can reveal low-dimensional structure.

Principal component analysis (PCA) is an orthogonal coordinate transform that is suitable for dimension reduction if some of the inputs are linearly correlated. In this case, PCA transforms redundant variables in the input space producing uncorrelated variables in the feature space.

There are a number of ways to derive the optimal PCA transform. One approach presented in [6] is based on finding uncorrelated features. It is straightforward to show that uncorrelated features have a diagonal covariance matrix. This can be used to solve for the covariance matrix C of input variables. By asserting the orthogonality of the PCA transform, we obtain V from the diagonalization of the covariance matrix $C = VDV^\top$. Given this PCA transform, we can show that V minimizes projection residuals as in [13].

2.1 Finding uncorrelated features

The correlation between two random variables x and y is defined as

$$\text{corr}(x, y) = \frac{E[(x - \mu_x)(y - \mu_y)]}{\sigma_x \sigma_y}, \quad (3)$$

where μ_x, μ_y and σ_x, σ_y are the respective means and standard deviations of x and y . We say x and y are uncorrelated when $\text{corr}(x, y) = 0$. This happens if and only if

$$E[(x - \mu_x)(y - \mu_y)] = \text{cov}(x, y) = 0. \quad (4)$$

The covariance matrix for a multivariate random variable $x = [x_1, x_2, \dots, x_d]$ (as a row vector) has $\text{cov}(x_i, x_j)$ in the i -th row and j -th column. Then

$$E[(x - \mu_x)^\top (x - \mu_x)] = \begin{bmatrix} \text{cov}(x_1, x_1) & \text{cov}(x_1, x_2) & \cdots & \text{cov}(x_1, x_d) \\ \text{cov}(x_2, x_1) & \text{cov}(x_2, x_2) & \cdots & \text{cov}(x_2, x_d) \\ \vdots & \vdots & \ddots & \vdots \\ \text{cov}(x_d, x_1) & \text{cov}(x_d, x_2) & \cdots & \text{cov}(x_d, x_d) \end{bmatrix}. \quad (5)$$

If x_1, x_2, \dots, x_d are pairwise uncorrelated, then $\text{cov}(x_i, x_j) = 0$ for all $i \neq j$. Hence, uncorrelated variables have a diagonal covariance matrix.

Now, let $a_1, a_2, \dots, a_n \in \mathbb{R}^{1 \times d}$ represent n observations in d variables. These observations can be considered points in d -dimensional space whose centroid is $\mu_a = \frac{1}{n} \sum_{i=1}^n a_i$. We want to determine a PCA transform which sends these points in the input space to points in the feature space. Moreover, the basis vectors of the feature space shall be uncorrelated. Accordingly, let $V \in \mathbb{R}^{d \times d}$ be the change of basis matrix and let

$$b_i = (a_i - \mu_a)V, \quad \text{for } i = 1, 2, \dots, n \quad (6)$$

be observations with respect to the feature coordinates. Then

$$\mu_b = \frac{1}{n} \sum_{i=1}^n b_i = \frac{1}{n} \sum_{i=1}^n (a_i - \mu_a) V = 0. \quad (7)$$

Using Equation (5), we can compute the sample covariance matrices as

$$C = \frac{1}{n-1} \sum_{i=1}^n (a_i - \mu_a)^\top (a_i - \mu_a), \quad D = \frac{1}{n-1} \sum_{i=1}^n b_i^\top b_i. \quad (8)$$

Since D is the covariance matrix of uncorrelated features, by the argument above, it is diagonal. If we restrict V to be orthogonal, then

$$b_i^\top b_i = V^\top (a_i - \mu_a)^\top (a_i - \mu_a) V \implies D = V^\top C V \implies C = V D V^\top. \quad (9)$$

Hence, V must be a matrix of orthonormal eigenvectors v_1, v_2, \dots, v_d corresponding to eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_d$ on the diagonal of D . When the eigenvalues and eigenvectors are ordered such that

$$\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d, \quad (10)$$

we call v_1, v_2, \dots, v_d the *principal components* of the PCA transform matrix V .

2.2 Minimizing projection residuals

Let $a_1, a_2, \dots, a_n \in \mathbb{R}^d$ be points in d -dimensions. Consider the orthonormal basis vectors $v_1, v_2, \dots, v_p \in \mathbb{R}^d$ for a subspace $F^p \subseteq \mathbb{R}^d$. The projection of these points onto F^p results in the points

$$b_i = \sum_{j=1}^p \langle a_i, v_j \rangle v_j, \quad \text{for } i = 1, 2, \dots, n. \quad (11)$$

The squared projection residuals are

$$\begin{aligned} \|a_i - b_i\|^2 &= \langle a_i - b_i, a_i - b_i \rangle \\ &= \langle a_i, a_i \rangle - 2 \langle a_i, b_i \rangle + \langle b_i, b_i \rangle \\ &= \|a_i\|^2 - 2 \langle a_i, b_i \rangle + \|b_i\|^2 \\ &= \|a_i\|^2 - 2 \left\langle a_i, \sum_{j=1}^p \langle a_i, v_j \rangle v_j \right\rangle + \left\| \sum_{j=1}^p \langle a_i, v_j \rangle v_j \right\|^2 \\ &= \|a_i\|^2 - 2 \sum_{j=1}^p \langle a_i, v_j \rangle \langle a_i, v_j \rangle + \sum_{j=1}^p \langle a_i, v_j \rangle^2 \|v_j\|^2 \\ &= \|a_i\|^2 - 2 \sum_{j=1}^p \langle a_i, v_j \rangle^2 + \sum_{j=1}^p \langle a_i, v_j \rangle^2 \\ &= \|a_i\|^2 - \sum_{j=1}^p \langle a_i, v_j \rangle^2 \end{aligned} \quad (12)$$

2.3 Singular value decomposition

2.4 Principal component analysis algorithm

Let A be a data matrix whose n rows correspond to observations and d columns correspond to variables. The following algorithm demonstrates a simple method for computing the PCA of A :

1. Compute the centered matrix $A_0 = A - \text{col mean}(A)$.
2. Compute the covariance matrix $C = \frac{1}{n-1} A_0^\top A_0$.
3. Diagonalize the covariance matrix such that $C = V D V^\top$.
4. Order the eigenvalues and eigenvectors so that $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$. We call the ordered eigenvalues the *principal components*.
5. Choose the dimension of the subspace $p \leq d$.
6. Construct the $d \times p$ projection matrix V_p using the first p principal components v_1, v_2, \dots, v_p .

Example 2.1. Consider the following matrix

$$A = \begin{bmatrix} 5 & 3 & 6 & 7 & 6 \\ 4 & 5 & 7 & 1 & 3 \\ 5 & 7 & 6 & 1 & 0 \\ 6 & 10 & 12 & 12 & 11 \\ 9 & 10 & 12 & 13 & 9 \end{bmatrix}.$$

The column means are $\mu = [5.8, 7, 8.6, 6.8, 5.8]$. Then the mean-centered data becomes

$$X = A - \mu = \frac{1}{5} \begin{bmatrix} -4 & -20 & -13 & 1 & 1 \\ -9 & -10 & -8 & -29 & -14 \\ -4 & 0 & -13 & -29 & -29 \\ 1 & 15 & 17 & 26 & 26 \\ 16 & 15 & 17 & 31 & 16 \end{bmatrix}.$$

The covariance matrix is

$$C = X^T X = \frac{1}{5} \begin{bmatrix} 74 & 85 & 93 & 179 & 104 \\ 85 & 190 & 170 & 225 & 150 \\ 93 & 170 & 196 & 313 & 238 \\ 179 & 225 & 313 & 664 & 484 \\ 104 & 150 & 238 & 484 & 394 \end{bmatrix}.$$

Diagonalizing C gives

$$V = \begin{bmatrix} 0.1888 & -0.2020 & -0.6366 & 0.5495 & -0.4651 \\ 0.2755 & -0.7886 & 0.1472 & -0.4502 & -0.2791 \\ 0.3606 & -0.3464 & 0.3128 & 0.5836 & 0.5582 \\ 0.6979 & 0.2522 & -0.4422 & -0.3707 & 0.3411 \\ 0.5209 & 0.3922 & 0.5288 & 0.1316 & -0.5271 \end{bmatrix},$$

$$D = \begin{bmatrix} 264.8458 & 0 & 0 & 0 & 0 \\ 0 & 27.9766 & 0 & 0 & 0 \\ 0 & 0 & 9.3198 & 0 & 0 \\ 0 & 0 & 0 & 1.4579 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}.$$

If we keep all 5 principal component vectors, then $V_5 = V$ and the projection of X along V is

$$P = XV = \begin{bmatrix} -1.9469 & 4.3453 & -0.8756 & -0.2039 & 0 \\ -6.9742 & -0.0660 & 1.4352 & 0.7590 & 0 \\ -8.1577 & -2.6752 & -0.8063 & -0.5704 & 0 \\ 8.4282 & -0.2330 & 1.8282 & -0.4996 & 0 \\ 8.6507 & -1.3711 & -1.5815 & 0.5149 & 0 \end{bmatrix}.$$

Here, the last column of P is the zero vector because the last eigenvalue of C is zero². To perfectly reconstruct A , we need $k = 4$ principal components and the row vector μ

$$A = PV^T + \mu = PV_4^T + \mu.$$

If we use $k = 3$ principal components, then the projection of X onto V_3 is

$$P = XV_3 = \begin{bmatrix} -1.9469 & 4.3453 & -0.8756 \\ -6.9742 & -0.0660 & 1.4352 \\ -8.1577 & -2.6752 & -0.8063 \\ 8.4282 & -0.2330 & 1.8282 \\ 8.6507 & -1.3711 & -1.5815 \end{bmatrix}$$

and A is approximately reconstructed by

$$A \approx PV_3^T + \mu = \begin{bmatrix} 5.1 & 2.9 & 6.1 & 6.9 & 6.0 \\ 3.6 & 5.3 & 6.6 & 1.3 & 2.9 \\ 5.3 & 6.7 & 6.3 & 0.8 & 0.1 \\ 6.3 & 9.8 & 12.3 & 11.8 & 11.1 \\ 8.7 & 10.2 & 11.7 & 13.2 & 8.9 \end{bmatrix}.$$

We can compute the reconstruction error using

$$E_k = \|A - (PV_k^T + \mu)\|_F,$$

²Since we subtracted the column means from a square matrix A , the dimension of the row space was reduced to 4.

where $\|\cdot\|_F$ is the Frobenius norm. By the SVD, we have $X = USV^T$, where $S = \sqrt{D}$. So, the projection of X onto V_k is

$$P = XV_k = US_k,$$

where S_k is the diagonal matrix of the first k singular values. Then the reconstruction error becomes

$$\begin{aligned} \|A - (PV_k^T + \mu)\|_F &= \|(A - \mu) - PV_k^T\|_F \\ &= \|X - PV_k^T\|_F \\ &= \|USV^T - US_kV^T\|_F \\ &= \|U(S - S_k)V^T\|_F \\ &= \|S - S_k\|_F \\ &= \sigma_k + \sigma_{k+1} + \cdots + \sigma_p. \end{aligned}$$

Hence,

$$E_3 = \sigma_3 + \sigma_4 = \sqrt{1.4579} + 0 = 1.2074.$$

2.5 Linear regression and PCA

[14] Let $\mathbf{x}_1, \dots, \mathbf{x}_m \in \mathbb{R}^n$ be observation vectors and $\mathbf{y} \in \mathbb{R}^n$ be a target vector. A linear regression model finds a vector of weights $\mathbf{w} = (w_1, w_2, \dots, w_n)$ to determine the linear function

$$f(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} = \sum_{i=1}^n w_i x_i, \quad (13)$$

where $\mathbf{x} = (x_1, x_2, \dots, x_n)$ is an input vector. The residual error of each observation is $\mathbf{y} - f(\mathbf{x}_i)$, for $i = 1, 2, \dots, m$. Then the residual sum of squares is given to be

$$RSS = \sum_{i=1}^m (\mathbf{y} - f(\mathbf{x}_i))^2 = \sum_{i=1}^m (\mathbf{y} - \mathbf{w} \cdot \mathbf{x}_i)^2 = (\mathbf{y} - X^\top \mathbf{w})^\top (\mathbf{y} - X^\top \mathbf{w}), \quad (14)$$

where $X = [\mathbf{x}_1, \dots, \mathbf{x}_m]$. By minimizing RSS , the magnitude of the residuals will be as small as possible, producing the optimal linear model $f(\mathbf{x})$. Therefore, this method is known as least squares regression. Differentiate (14) with respect to \mathbf{w} and set equal to zero so that

$$2X^\top \mathbf{y} - 2X^\top X \mathbf{w} = 0. \quad (15)$$

This leads to the normal equation $X^\top X \mathbf{w} = X^\top \mathbf{y}$. Thus,

$$\mathbf{w} = (X^\top X)^{-1} X^\top \mathbf{y} \quad (16)$$

gives the optimal linear model which minimizes residual error.

Example 2.2. In two dimensions, let $\mathbf{x} = (x_1, \dots, x_m)$ and $\mathbf{y} = (y_1, \dots, y_m)$. Then the least squares regression model is given by

$$f(x) = \frac{\text{cov}(\mathbf{x}, \mathbf{y})}{\text{var}(\mathbf{x})}(x - \bar{\mathbf{x}}) + \bar{\mathbf{y}}, \quad (17)$$

Since \mathbf{x} is treated as an input variable and \mathbf{y} is treated as an output, this is a type of supervised learning. In contrast, PCA is an unsupervised learning technique. PCA organizes the variables in the input space to reveal any patterns in the underlying data. If we have input variables \mathbf{x}_1 and \mathbf{x}_2 , we can use PCA to find a linear pattern among the inputs without trying to predict an output. First, we find the direction of the largest variance λ_{\max} using the covariance matrix

$$\begin{bmatrix} \text{cov}(\mathbf{x}_1, \mathbf{x}_1) & \text{cov}(\mathbf{x}_1, \mathbf{x}_2) \\ \text{cov}(\mathbf{x}_2, \mathbf{x}_1) & \text{cov}(\mathbf{x}_2, \mathbf{x}_2) \end{bmatrix} = \begin{bmatrix} \text{var}(\mathbf{x}_1) & \text{cov}(\mathbf{x}_1, \mathbf{x}_2) \\ \text{cov}(\mathbf{x}_1, \mathbf{x}_2) & \text{var}(\mathbf{x}_2) \end{bmatrix}.$$

In particular,

$$\lambda_{\max} = \frac{1}{2} \left(\text{var}(\mathbf{x}) + \text{var}(\mathbf{y}) + \sqrt{(\text{var}(\mathbf{x}) - \text{var}(\mathbf{y}))^2 + 4 \text{cov}(\mathbf{x}, \mathbf{y})^2} \right).$$

Then the PCA regression model becomes

$$g(x) = \frac{\lambda_{\max} - \text{var}(\mathbf{x}_1)}{\text{cov}(\mathbf{x}_1, \mathbf{x}_2)}(x - \bar{\mathbf{x}}_1) + \bar{\mathbf{x}}_2. \quad (18)$$

In this case, the projection residuals are orthogonal to the PCA regression line. See Figure 2.

3 Reproducing Kernel Hilbert Space

In this section, our goal is to establish properties of Hilbert spaces and kernel functions that can be used to modify the PCA algorithm. To begin, we will briefly cite some definitions and results from analysis [7], [10] and matrix theory [5].

Definition 3.1. [15] Let X be a (real) vector space. An *inner product* is a function $\langle \cdot, \cdot \rangle : X \times X \rightarrow \mathbb{R}$ which satisfies the following properties:

1. Symmetry. For all $x, y \in X$,

$$\langle x, y \rangle = \langle y, x \rangle.$$

2. Linear in the first argument. For all $x, y, z \in X$, $\alpha, \beta \in \mathbb{R}$,

$$\langle \alpha x + \beta y, z \rangle = \alpha \langle x, z \rangle + \beta \langle y, z \rangle.$$

3. Positive definite. For all $x \in X$,

$$\langle x, x \rangle \geq 0$$

and $\langle x, x \rangle = 0$ if and only if $x = 0$.

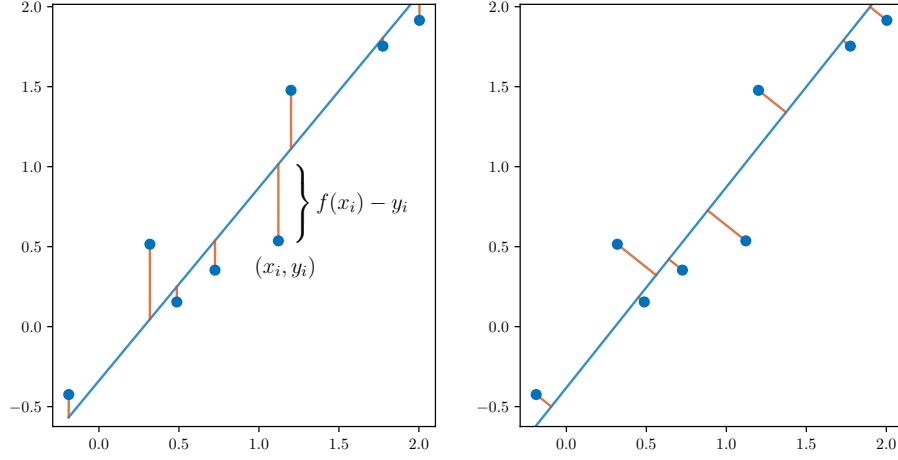


Figure 2: Least squares regression model $f(x)$ (left) minimizes the sum of squared errors while total least squares, i.e., the PCA model $g(x)$ (right) minimizes the orthogonal projections.

An *inner product space* is a vector space along with an inner product.

Since real inner products are symmetric and linear in the first argument,

$$\langle z, \alpha x + \beta y \rangle = \langle \alpha x + \beta y, z \rangle = \alpha \langle x, z \rangle + \beta \langle y, z \rangle = \alpha \langle z, x \rangle + \beta \langle z, y \rangle.$$

So, we also have linearity in the second argument.

The norm induced by an inner product is defined as

$$\|x\| = \langle x, x \rangle^{1/2}$$

and the metric induced by this norm is

$$d(x, y) = \|x - y\| = \langle x - y, x - y \rangle^{1/2}.$$

It follows that an inner product space is also a normed space and a metric space. So, the induced norm will have the following properties for all $x, y \in X$ and $\alpha \in \mathbb{R}$:

1. Triangle inequality. $\|x + y\| \leq \|x\| + \|y\|$;
2. Scalar multiplication. $\|\alpha x\| = |\alpha| \|x\|$;
3. Positivity. $\|x\| \geq 0$ and $\|x\| = 0$ if and only if $x = 0$.

Definition 3.2 (Hilbert space). A Hilbert space is a complete inner product space. For a Hilbert space H , we sometimes denote the inner product as $\langle \cdot, \cdot \rangle_H$ to avoid ambiguity.

Two Hilbert spaces H and L (over the same field) are said to be *isomorphic* if there is a bijection $T : H \rightarrow L$ such that

$$\langle x, y \rangle_H = \langle Tx, Ty \rangle_L, \quad (19)$$

for every $x, y \in H$. In [7], Kreyszig shows that two Hilbert spaces are isomorphic if and only if they have the same dimension. If V is an inner product space that is not complete, then it can be extended to a Hilbert space by completion. The completion of an inner product space is denoted as \overline{V} and is unique up to isomorphism.

A Hilbert space is said to be *separable* if it contains a dense countable subset. It can be shown [7] that a Hilbert space is separable if and only if it has a countable orthonormal basis. The following example demonstrates a useful property of separable Hilbert spaces.

Example 3.3. [10] The space of square-summable (real) sequences is defined as

$$\ell^2(A) = \left\{ x : A \rightarrow \mathbb{R} \mid \sum_{a \in A} x_a^2 < \infty \right\}. \quad (20)$$

Given the inner product

$$\langle x, y \rangle = \sum_{a \in A} x_a y_a, \quad (21)$$

$\ell^2(A)$ is a Hilbert space. Moreover, $\ell^2(A)$ is separable if and only if A is countable. It follows that the sequence space $\ell^2 = \ell^2(\mathbb{N})$ is the separable Hilbert space of square-summable sequences. Due to the Riesz-Fischer theorem, every infinite-dimensional Hilbert space is isomorphic to ℓ^2 .

Definition 3.4 (Gram matrix). [5] Let $x_1, x_2, \dots, x_n \in X$ for some inner product space X equipped with $\langle \cdot, \cdot \rangle$. We say G is a *Gram matrix* (or *Gramian*) for the set of vectors $\{x_1, x_2, \dots, x_n\}$ with respect to $\langle \cdot, \cdot \rangle$ if $G = [\langle x_i, x_j \rangle]_{ij}$.

Example 3.5. Consider the vectors in \mathbb{R}^3 :

$$\mathbf{v}_1 = \begin{bmatrix} v_{11} \\ v_{21} \\ v_{31} \end{bmatrix}, \quad \mathbf{v}_2 = \begin{bmatrix} v_{12} \\ v_{22} \\ v_{32} \end{bmatrix}, \quad \mathbf{v}_3 = \begin{bmatrix} v_{13} \\ v_{23} \\ v_{33} \end{bmatrix}, \quad \mathbf{v}_4 = \begin{bmatrix} v_{14} \\ v_{24} \\ v_{34} \end{bmatrix}.$$

The Gram matrix for these vectors is

$$G = \begin{bmatrix} \mathbf{v}_1^\top \mathbf{v}_1 & \mathbf{v}_1^\top \mathbf{v}_2 & \mathbf{v}_1^\top \mathbf{v}_3 & \mathbf{v}_1^\top \mathbf{v}_4 \\ \mathbf{v}_2^\top \mathbf{v}_1 & \mathbf{v}_2^\top \mathbf{v}_2 & \mathbf{v}_2^\top \mathbf{v}_3 & \mathbf{v}_2^\top \mathbf{v}_4 \\ \mathbf{v}_3^\top \mathbf{v}_1 & \mathbf{v}_3^\top \mathbf{v}_2 & \mathbf{v}_3^\top \mathbf{v}_3 & \mathbf{v}_3^\top \mathbf{v}_4 \\ \mathbf{v}_4^\top \mathbf{v}_1 & \mathbf{v}_4^\top \mathbf{v}_2 & \mathbf{v}_4^\top \mathbf{v}_3 & \mathbf{v}_4^\top \mathbf{v}_4 \end{bmatrix}.$$

If V is a matrix whose columns are $\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3, \mathbf{v}_4$, then we can write $G = V^\top V$.

Theorem 3.6. [5] *A matrix G is a Gram matrix if and only if G is symmetric and positive semidefinite.*

Proof. (\Rightarrow) Suppose G is the Gram matrix of x_1, x_2, \dots, x_n with respect to $\langle \cdot, \cdot \rangle$. Then G is symmetric because $G_{ij} = \langle x_i, x_j \rangle = \langle x_j, x_i \rangle = G_{ji}$.

Let $c_1, c_2, \dots, c_n \in \mathbb{R}$. Then G is positive semidefinite because

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j \langle x_i, x_j \rangle = \left\langle \sum_{i=1}^n c_i x_i, \sum_{j=1}^n c_j x_j \right\rangle = \left\| \sum_{i=1}^n c_i x_i \right\|^2 \geq 0. \quad (22)$$

(\Leftarrow) Suppose G is symmetric and positive semidefinite. Then G can be diagonalized as $G = V^\top D V = (D^{1/2} V)^\top (D^{1/2} V)$. Let $w_i = \sqrt{\lambda_i} v_i$, for $i = 1, 2, \dots, n$, where λ_i and v_i are eigenvalues and eigenvectors of G , respectively. Then $[G]_{ij} = w_i^\top w_j$. Hence G is a Gram matrix. \square

Theorem 3.7. [5] *A Gram matrix G of x_1, x_2, \dots, x_n is positive definite if and only if x_1, x_2, \dots, x_n are linearly independent.*

Proof. \square

Definition 3.8. A *symmetric bilinear form* is a map $k : X \times X \rightarrow \mathbb{R}$ over a vector space X such that, for all $x, y, z \in X$, $\alpha, \beta \in \mathbb{R}$,

1. $k(x, y) = k(y, x)$ and
2. $k(\alpha x + \beta y, z) = \alpha k(x, z) + \beta k(y, z)$.

This can be thought of as a generalization of an inner product which is symmetric and bilinear, but not necessarily positive definite. If $U = \{u_1, u_2, \dots, u_n\}$ is a basis for X , then we can define a matrix $A = [k(u_i, u_j)]_{ij}$. Clearly, A is symmetric since $k(u_i, u_j) = k(u_j, u_i)$. Let $v = \sum_{i=1}^n \alpha_i u_i$ and $w = \sum_{i=1}^n \beta_i u_i$ be vectors with respect to the basis U and let $x = [\alpha_i]_{i=1}^n$ and $y = [\beta_i]_{i=1}^n$. Then we can write

$$k(v, w) = k \left(\sum_{i=1}^n \alpha_i u_i, \sum_{j=1}^n \beta_j u_j \right) = \sum_{i=1}^n \sum_{j=1}^n \alpha_i \beta_j k(u_i, u_j) = x^\top A y. \quad (23)$$

If $A = I$, then $v = x$, $w = y$, and $k(v, w) = v^\top w$ is simply the dot product.

Using Equation (23), we say that k is positive definite whenever A is positive definite. By Theorem 3.6, A is a Gram matrix. It follows that k is associated with some inner product space.

If X is an infinite vector space, we cannot construct the basis matrix A . In practice, we can use a subspace spanned by a subset of vectors $x_1, x_2, \dots, x_n \in X$ to generate A for some symmetric bilinear map k .

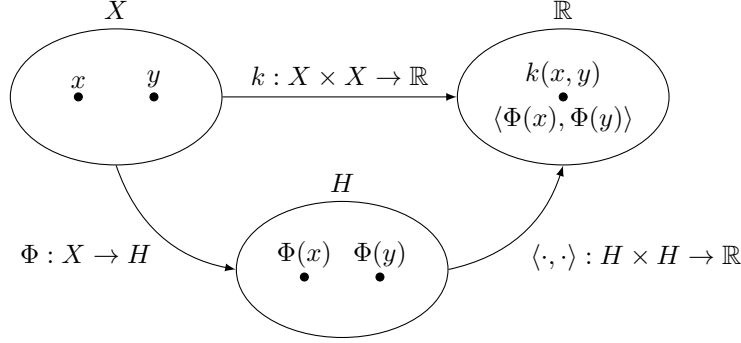


Figure 3: Kernel map diagram.

3.1 Kernel methods

The development of *kernel functions* can be traced back to the beginning of the twentieth century when David Hilbert and James Mercer were studying integral equations [4]. Hilbert proved some important results in [3] about the eigenvalues of an integral operator whose kernel function is of *definite* type. Expanding on Hilbert's work, Mercer provided the necessary conditions in [8] that allow a kernel function to be written in terms of the eigenvalues and eigenfunctions of the integral operator. This result became known as Mercer's theorem. See Appendix C. As a corollary, a kernel function which satisfies Mercer's conditions can be written as an inner product.

Hilbert spaces and Mercer's theorem led to a number of advances in functional analysis. In 1950, Nachman Aronszajn introduced *reproducing kernels* and their associated Hilbert spaces [1]. Later, the work of Mercer and Aronszajn inspired the application of kernels in machine learning. So-called kernel methods are ways to adapt a machine learning algorithm by replacing a dot product with a kernel function. In Section 4, we will look at the kernel method applied to the PCA algorithm. For now, we will examine the mathematics behind kernel methods.

Let X be a nonempty set and $k : X \times X \rightarrow \mathbb{R}$ be a kernel function. We want k to coincide with the inner product of some Hilbert space H . Then the elements of X should be associated with the elements of H via a feature map $\Phi : X \rightarrow H$ so that

$$k(x, y) = \langle \Phi(x), \Phi(y) \rangle, \quad (24)$$

for all $x, y \in X$. We can think of the kernel k as a shortcut which allows us to bypass the potentially high-dimensional space H . See Figure 3.

Now, we want to restrict k so that it is a valid inner product. Since inner products are symmetric and positive definite, it follows that k should have these properties as well. Let $S = \{x_1, x_2, \dots, x_n\} \subseteq X$ be a sample of points in X . We say K is the *kernel matrix* of S with respect to k if $[K]_{ij} = k(x_i, x_j)$ for all $i, j = 1, 2, \dots, n$. Then K should be equal to the Gram matrix of $\Phi(S)$ with

respect to $\langle \cdot, \cdot \rangle$. By Theorem 3.6, K will be symmetric and positive semidefinite.

Definition 3.9 (kernel). Let X be a nonempty set and $k : X \times X \rightarrow \mathbb{R}$. Then k is a (positive definite) *kernel* if:

1. k is symmetric. For all $x, y \in X$, $k(x, y) = k(y, x)$.
2. k is positive definite: if $x_1, \dots, x_n \in X$ and $c_1, \dots, c_n \in \mathbb{R}$, then

$$\sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) \geq 0.$$

Equivalently, the kernel matrix $K \in \mathbb{R}^{n \times n}$ whose entries are $[K]_{ij} = k(x_i, x_j)$ is positive semidefinite, that is, $\mathbf{c}^\top K \mathbf{c} \geq 0$ for all $\mathbf{c} \in \mathbb{R}^n$.

3.2 Constructing kernels

Theorem 3.10. [9, 14] Suppose k_1 and k_2 are kernels over $X \times X$. The following functions kernels.

1. $k(x, y) = a_1 k_1(x, y) + a_2 k_2(x, y)$ for all $a_1, a_2 \geq 0$.
2. $k(x, y) = k_1(x, y) k_2(x, y)$.
3. $k(x, y) = a_0 + a_1 k_1(x, y) + a_2 k_1(x, y)^2 + \dots + a_n k_1(x, y)^n$ for all $n \in \mathbb{N}$ and $a_0, \dots, a_n \geq 0$.
4. $k(x, y) = k_1(h(x), h(y))$ for all $h : X \rightarrow X$.
5. $k(x, y) = g(x)g(y)$ for all $g : X \rightarrow \mathbb{R}$.
6. $k(x, y) = \exp(k_1(x, y))$.

Proof. Let $x_1, \dots, x_n \in X$ and $c_1, \dots, c_n \in \mathbb{R}$.

1. Let $k = a_1 k_1 + a_2 k_2$ for $a_1, a_2 \geq 0$. Since k_1 and k_2 are symmetric,

$$k(x, y) = a_1 k_1(x, y) + a_2 k_2(x, y) = a_1 k_1(y, x) + a_2 k_2(y, x) = k(y, x),$$

for all $x, y \in X$. So, k is symmetric.

Since k_1 and k_2 are positive semidefinite and $a_1, a_2 \geq 0$,

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j (a_1 k_1(x_i, x_j) + a_2 k_2(x_i, x_j)) \\ &= a_1 \sum_{i=1}^n \sum_{j=1}^n c_i c_j k_1(x_i, x_j) + a_2 \sum_{i=1}^n \sum_{j=1}^n c_i c_j k_2(x_i, x_j) \\ &\geq 0. \end{aligned}$$

So, k is positive semidefinite.

2. Let $k = k_1 k_2$. Define K so that $[K]_{ij} = k(x_i, x_j) = k_1(x_i, x_j)k_2(x_i, x_j)$. Let K_1 and K_2 be the Gram matrices for k_1 and k_2 , respectively. Then K_1, K_2 have orthonormal eigenvectors and nonnegative eigenvalues such that

$$\begin{aligned}
K_1 &= VLV^\top \\
&= \begin{bmatrix} v_{11} & \cdots & v_{1n} \\ \vdots & \ddots & \vdots \\ v_{n1} & \cdots & v_{nn} \end{bmatrix} \begin{bmatrix} \lambda_1 & \cdots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \cdots & \lambda_n \end{bmatrix} \begin{bmatrix} v_{11} & \cdots & v_{1n} \\ \vdots & \ddots & \vdots \\ v_{1n} & \cdots & v_{nn} \end{bmatrix} \\
&= \begin{bmatrix} \sum_{j=1}^n \lambda_j v_{1j} v_{1j} & \cdots & \sum_{j=1}^n \lambda_j v_{nj} v_{1j} \\ \vdots & \ddots & \vdots \\ \sum_{j=1}^n \lambda_j v_{1j} v_{nj} & \cdots & \sum_{j=1}^n \lambda_j v_{nj} v_{nj} \end{bmatrix} \\
&= \sum_{j=1}^n \lambda_j \begin{bmatrix} v_{1j} v_{1j} & \cdots & v_{nj} v_{1j} \\ \vdots & \ddots & \vdots \\ v_{1j} v_{nj} & \cdots & v_{nj} v_{nj} \end{bmatrix}
\end{aligned}$$

and

$$K_2 = UMU^\top = \sum_{j=1}^n \mu_j \begin{bmatrix} u_{1j} u_{1j} & \cdots & u_{nj} u_{1j} \\ \vdots & \ddots & \vdots \\ u_{1j} u_{nj} & \cdots & u_{nj} u_{nj} \end{bmatrix}.$$

Let $\mathbf{v}_i = [v_{1i} \ \cdots \ v_{ni}]^\top$ and $\mathbf{u}_j = [u_{1j} \ \cdots \ u_{nj}]$, for all $i, j = 1, 2, \dots, n$. Then

$$\begin{aligned}
K &= K_1 \circ K_2 \\
&= \sum_{i=1}^n \lambda_i \begin{bmatrix} v_{1i} v_{1i} & \cdots & v_{ni} v_{1i} \\ \vdots & \ddots & \vdots \\ v_{1i} v_{ni} & \cdots & v_{ni} v_{ni} \end{bmatrix} \circ \sum_{j=1}^n \mu_j \begin{bmatrix} u_{1j} u_{1j} & \cdots & u_{nj} u_{1j} \\ \vdots & \ddots & \vdots \\ u_{1j} u_{nj} & \cdots & u_{nj} u_{nj} \end{bmatrix} \\
&= \sum_{i=1}^n \sum_{j=1}^n \lambda_i \mu_j \begin{bmatrix} v_{1i} u_{1j} v_{1i} u_{1j} & \cdots & v_{1i} u_{1j} v_{ni} u_{nj} \\ \vdots & \ddots & \vdots \\ v_{ni} u_{nj} v_{1i} u_{1j} & \cdots & v_{ni} u_{nj} v_{ni} u_{nj} \end{bmatrix} \\
&= \sum_{i=1}^n \sum_{j=1}^n \lambda_i \mu_j \begin{bmatrix} v_{1i} u_{1j} \\ \vdots \\ v_{ni} u_{nj} \end{bmatrix} [v_{1i} u_{1j} \ \cdots \ v_{ni} u_{nj}] \\
&= \sum_{i=1}^n \sum_{j=1}^n \lambda_i \mu_j (\mathbf{v}_i \circ \mathbf{u}_j) (\mathbf{v}_i \circ \mathbf{u}_j)^\top,
\end{aligned}$$

where \circ is the Hadamard product. Each $(\mathbf{v}_i \circ \mathbf{u}_j)(\mathbf{v}_i \circ \mathbf{u}_j)^\top$ is a symmetric positive semidefinite matrix. Since K_1, K_2 are positive semidefinite, we have $\lambda_i, \mu_i > 0$. Then K is symmetric positive semidefinite.

3. By part 2, k_1, k_1^2, \dots, k_1^n are kernels. By part 1, $a_0 + a_1 k_1 + a_2 k_1^2 + \dots + a_n k_1^n$ is a kernel.
4. Since $y_i = h(x_i) \in X$ for all $i = 1, 2, \dots, n$, we have

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j k_1(h(x_i), h(x_j)) \\ &= \sum_{i=1}^n \sum_{j=1}^n c_i c_j k_1(y_i, y_j) \\ &\geq 0. \end{aligned}$$

5. Let $g : X \rightarrow \mathbb{R}$ and let $c_i g(x_i) = y_i \in \mathbb{R}$. If $k(x, y) = g(x)g(y)$, then

$$\begin{aligned} \sum_{i=1}^n \sum_{j=1}^n c_i c_j k(x_i, x_j) &= \sum_{i=1}^n \sum_{j=1}^n c_i g(x_i) c_j g(x_j) \\ &= \sum_{i=1}^n \sum_{j=1}^n y_i y_j \\ &= \left(\sum_{i=1}^n y_i \right)^2 \\ &\geq 0. \end{aligned}$$

6. Let K_1 be the Gram matrix for k_1 . If $K_1 v = \lambda v$, then $K_1^m v = \lambda^m v$ for all $m \in \mathbb{N}$. So,

$$(\exp K_1)v = \sum_{m=0}^{\infty} \frac{K_1^m v}{m!} = \sum_{m=0}^{\infty} \frac{\lambda^m v}{m!} = e^\lambda v.$$

Then $K = \exp K_1$ has eigenvalues e^λ . Since K_1 is positive semidefinite, it has real eigenvalues so that $e^\lambda > 0$. It follows that K is positive definite.

□

Theorem 3.11 (Gaussian kernel). *The function $k : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ defined by*

$$k(x, y) = \exp\left(\frac{-\|x - y\|_2^2}{\sigma^2}\right),$$

is a kernel.

Proof.

□

4 Kernel PCA

Recall that linear PCA finds new components that reveal more information about the structure of high-dimensional data. Since PCA is an orthogonal projection, the original data is rotated within the original space of input variables. The work of Schölkopf, Smola, and Müller [11, 12] generalized PCA based on the successful application of kernel methods in support vector machines. In kernel PCA, the inner product of the input space is replaced with the inner product of a feature space. As such, the principal components of kernel PCA are nonlinear transformations of input variables, or features.

Using results from the previous section, a kernel function k defines a unique reproducing kernel Hilbert space H_k and feature map $\Phi : X \rightarrow H_k$ such that

$$k(x, y) = \langle \Phi(x), \Phi(y) \rangle, \quad (25)$$

for all $x, y \in H_k$.

4.1 Covariance matrix and kernel matrix

In PCA, we diagonalize $X^\top X$.

4.2 Centering in the feature space

Let $\Phi : X \rightarrow H_k$ be a feature map determined by a kernel k . Since Φ may be nonlinear, the image $\Phi(x)$ of a centered vector $x \in X$ is not guaranteed to be centered. For an effective PCA algorithm, it is necessary to compute the kernel matrix of centered vectors in the feature space. [12]

Given $x_1, \dots, x_n \in X$, the points

$$\Phi_0(x_i) = \Phi(x_i) - \frac{1}{n} \sum_{i=1}^n \Phi(x_i), \quad \text{for } i = 1, \dots, n \quad (26)$$

are the centered feature vectors in H_k . Then the centered kernel matrix becomes

$$\begin{aligned}
[K_0]_{ij} &= \langle \Phi_0(x_i), \Phi_0(x_j) \rangle \\
&= \left\langle \Phi(x_i) - \frac{1}{n} \sum_{p=1}^n \Phi(x_p), \Phi(x_j) - \frac{1}{n} \sum_{q=1}^n \Phi(x_q) \right\rangle \\
&= \langle \Phi(x_i), \Phi(x_j) \rangle - \frac{1}{n} \sum_{p=1}^n \langle \Phi(x_p), \Phi(x_j) \rangle \\
&\quad - \frac{1}{n} \sum_{q=1}^n \langle \Phi(x_i), \Phi(x_q) \rangle \\
&\quad + \frac{1}{n^2} \sum_{p=1}^n \sum_{q=1}^n \langle \Phi(x_p), \Phi(x_q) \rangle \\
&= [K]_{ij} - \frac{1}{n} \sum_{p=1}^n [K]_{pj} - \frac{1}{n} \sum_{q=1}^n [K]_{iq} + \frac{1}{n^2} \sum_{p=1}^n \sum_{q=1}^n [K]_{pq},
\end{aligned}$$

where K is the uncentered kernel matrix given by $[K]_{ij} = \langle \Phi(x_i), \Phi(x_j) \rangle$. Then the formula for the centered kernel matrix can be written as

$$K_0 = K - \text{col mean}(K) - \text{row mean}(K) + \text{mean}(K). \quad (27)$$

See notes in Appendices A.1 and A.2.

Algorithm 1: Kernel PCA
<p>Input: Data matrix $A \in \mathbb{R}^{n \times d}$, kernel function k, number of components $p \leq n$</p> <p>Output: Transformed data $\tilde{A} \in \mathbb{R}^{n \times p}$</p>

Example 4.1. Consider the problem of classifying points based on their radii. These points cannot be separated using a linear classifier in the two dimensions. However, by mapping them to a three-dimensional space, they can be separated by planes. Applying kernel PCA, these points can be sent to the RKHS associated with a Gaussian kernel without using an explicit feature map. The points in this high-dimensional feature space can then be projected onto the first three principal components to find separation boundaries. See Figure 4.

5 Conclusion

A Linear Algebra

A number of matrix definitions and results are presented without proof. Unless otherwise specified, let A be an $n \times d$ matrix over the real numbers.

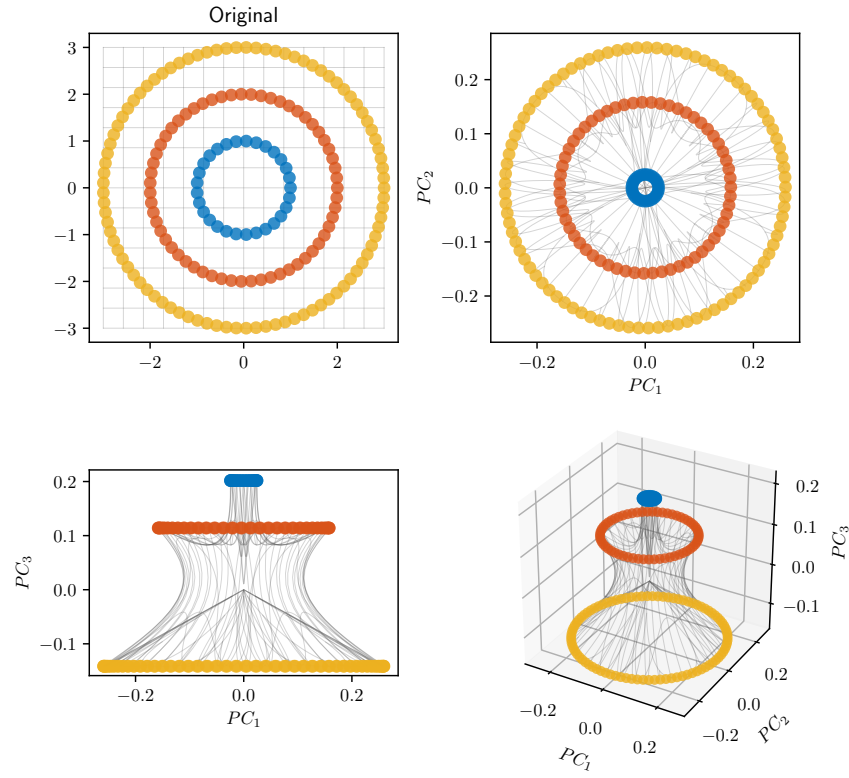


Figure 4: An idealized set of points in the plane are classified based on their radius. Kernel PCA with the Gaussian kernel is applied to find separation boundaries using the first three principal components.

1. A is *normal* if it is square and $AA^\top = A^\top A$.
2. Symmetric matrices are normal.
3. Symmetric matrices have orthogonal eigenvectors and real eigenvalues.
4. Positive semi-definite matrices have nonnegative eigenvalues.
5. Positive definite matrices have positive eigenvalues.
6. $A^\top A$ and AA^\top are symmetric positive semi-definite.

A.1 Matrix operations and notation

Let A be an $n \times d$ matrix. We write $[A]_{ij}$ to indicate the matrix entry in the i -th row and the j -th column.

Definition A.1. Define the *entry-wise mean* of A as

$$\text{mean}(A) = \frac{1}{nd} \sum_{i=1}^n \sum_{j=1}^d [A]_{ij}. \quad (28)$$

Define the *column-wise mean* of A as a $1 \times d$ row vector whose j -th entry is the mean of column j given by the formula

$$\begin{aligned} \text{col mean}(A) &= \left[\frac{1}{n} \sum_{i=1}^n [A]_{i1}, \quad \frac{1}{n} \sum_{i=1}^n [A]_{i2}, \quad \dots, \quad \frac{1}{n} \sum_{i=1}^n [A]_{id} \right] \\ &= \frac{1}{n} \sum_{i=1}^n \left[[A]_{i1}, \quad [A]_{i2}, \quad \dots, \quad [A]_{id} \right]. \end{aligned} \quad (29)$$

Define the *row-wise mean* of A as an $n \times 1$ column vector whose i -th entry is the mean of row i given by the formula

$$\text{row mean}(A) = \begin{bmatrix} \frac{1}{d} \sum_{j=1}^d [A]_{1j} \\ \frac{1}{d} \sum_{j=1}^d [A]_{2j} \\ \vdots \\ \frac{1}{d} \sum_{j=1}^d [A]_{nj} \end{bmatrix} = \frac{1}{d} \sum_{j=1}^d \begin{bmatrix} [A]_{1j} \\ [A]_{2j} \\ \vdots \\ [A]_{nj} \end{bmatrix}. \quad (30)$$

Let $[a]_{p \times q}$ denote the $p \times q$ repeated matrix whose entries are all a . Then Equations (28) to (30) can be written as

$$\text{mean}(A) = \left[\frac{1}{n} \right]_{1 \times n} \cdot A \cdot \left[\frac{1}{d} \right]_{d \times 1} \quad (31)$$

$$\text{col mean}(A) = \left[\frac{1}{n} \right]_{1 \times d} \cdot A \quad (32)$$

$$\text{row mean}(A) = A \cdot \left[\frac{1}{d} \right]_{d \times 1}. \quad (33)$$

A.2 Broadcasting

Consider the sum of two real matrices $A + B$. By definition, A and B must both have size $n \times d$. This means we cannot add a 2×3 matrix A and a 2×1 vector \mathbf{b} . However, in many programming languages the sum $A + \mathbf{b}$ would be handled using *broadcasting* [2]. In this case, \mathbf{b} is converted to a 2×3 matrix $\begin{bmatrix} \mathbf{b} & \mathbf{b} & \mathbf{b} \end{bmatrix}$ so that normal matrix addition applies. Generally, broadcasting a vector $\mathbf{b} \in \mathbb{R}^n$ to an $n \times d$ matrix can be represented as the matrix product

$$\mathbf{b} \cdot [1]_{1 \times d} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix} \cdot \begin{bmatrix} 1 & 1 & \cdots & 1 \end{bmatrix} = \begin{bmatrix} b_1 & b_1 & \cdots & b_1 \\ b_2 & b_2 & \cdots & b_2 \\ \vdots & \vdots & & \vdots \\ b_n & b_n & \cdots & b_n \end{bmatrix}, \quad (34)$$

where the notation $[a]_{n \times d}$ represents an $n \times d$ matrix whose entries are all a .

Definition A.2. For an $n \times d$ matrix A , we can define addition by an $n \times 1$ column vector \mathbf{c} as

$$A + \mathbf{c} := A + \mathbf{c} \cdot [1]_{1 \times d}. \quad (35)$$

Similarly, addition by a $1 \times d$ row vector \mathbf{r} can be defined as

$$A + \mathbf{r} := A + [1]_{n \times 1} \cdot \mathbf{r} \quad (36)$$

and addition by a scalar a can be defined as

$$A + a := A + a \cdot [1]_{n \times d}. \quad (37)$$

The left hand sides of Equations (35) to (37) are more concise and intuitive than the right hand sides. Provided that the vector types are clearly defined and compatible, there should be no ambiguity when adding column vectors, row vectors, and scalars to matrices. Moreover, this method of broadcasting is consistent with scientific programming languages.

B Riesz Representation Theorem

Theorem B.1 (Riesz Representation Theorem). [15] *Let $\phi : H \rightarrow \mathbb{R}$ be a continuous linear functional defined on a Hilbert space H . Then there exists a unique element $g \in H$ such that $\phi(g) = \langle f, g \rangle_H$ for all $g \in H$.*

Proof.

□

C Mercer’s Theorem

D Code

References

- [1] Nachman Aronszajn. Theory of reproducing kernels. *Transactions of the American mathematical society*, 68(3):337–404, 1950.
- [2] Charles R Harris, K Jarrod Millman, Stéfan J Van Der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J Smith, et al. Array programming with numpy. *Nature*, 585(7825):357–362, 2020.
- [3] D. Hilbert. *Grundzüge einer allgemeinen Theorie der linearen Integralgleichungen*. Cornell University Library historical math monographs. B. G. Teubner, 1912.
- [4] Thomas Hofmann, Bernhard Schölkopf, and Alexander J. Smola. Kernel methods in machine learning. *The Annals of Statistics*, 36(3), jun 2008.
- [5] R.A. Horn and C.R. Johnson. *Matrix Analysis*. Cambridge University Press, 2013.
- [6] K. Koutroumbas and S. Theodoridis. *Pattern Recognition*. Elsevier Science, 2008.
- [7] E. Kreyszig. *Introductory Functional Analysis with Applications*. Wiley Classics Library. Wiley, 1991.
- [8] James Mercer. Xvi. functions of positive and negative type, and their connection the theory of integral equations. *Philosophical transactions of the royal society of London. Series A, containing papers of a mathematical or physical character*, 209(441-458):415–446, 1909.
- [9] Cynthia Rudin. Intuition for the Algorithms of Machine Learning. Self-pub, 2020.
- [10] Walter Rudin. *Real and Complex Analysis*. Higher Mathematics Series. McGraw-Hill Education, 1987.
- [11] Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Kernel principal component analysis. In *International conference on artificial neural networks*, pages 583–588. Springer, 1997.
- [12] Bernhard Schölkopf, Alexander Smola, and Klaus-Robert Müller. Nonlinear component analysis as a kernel eigenvalue problem. *Neural computation*, 10(5):1299–1319, 1998.

- [13] Cosma Rohilla Shalizi. Advanced Data Analysis from an Elementary Point of View. Draft textbook, 2021.
- [14] John Shawe-Taylor and Nello Cristianini. *Kernel Methods for Pattern Analysis*. Cambridge University Press, 2004.
- [15] Christopher G. Small and D.L. Mcleish. *Hilbert Space Methods in Probability and Statistical Inference*. John Wiley & Sons, Inc, 1994.